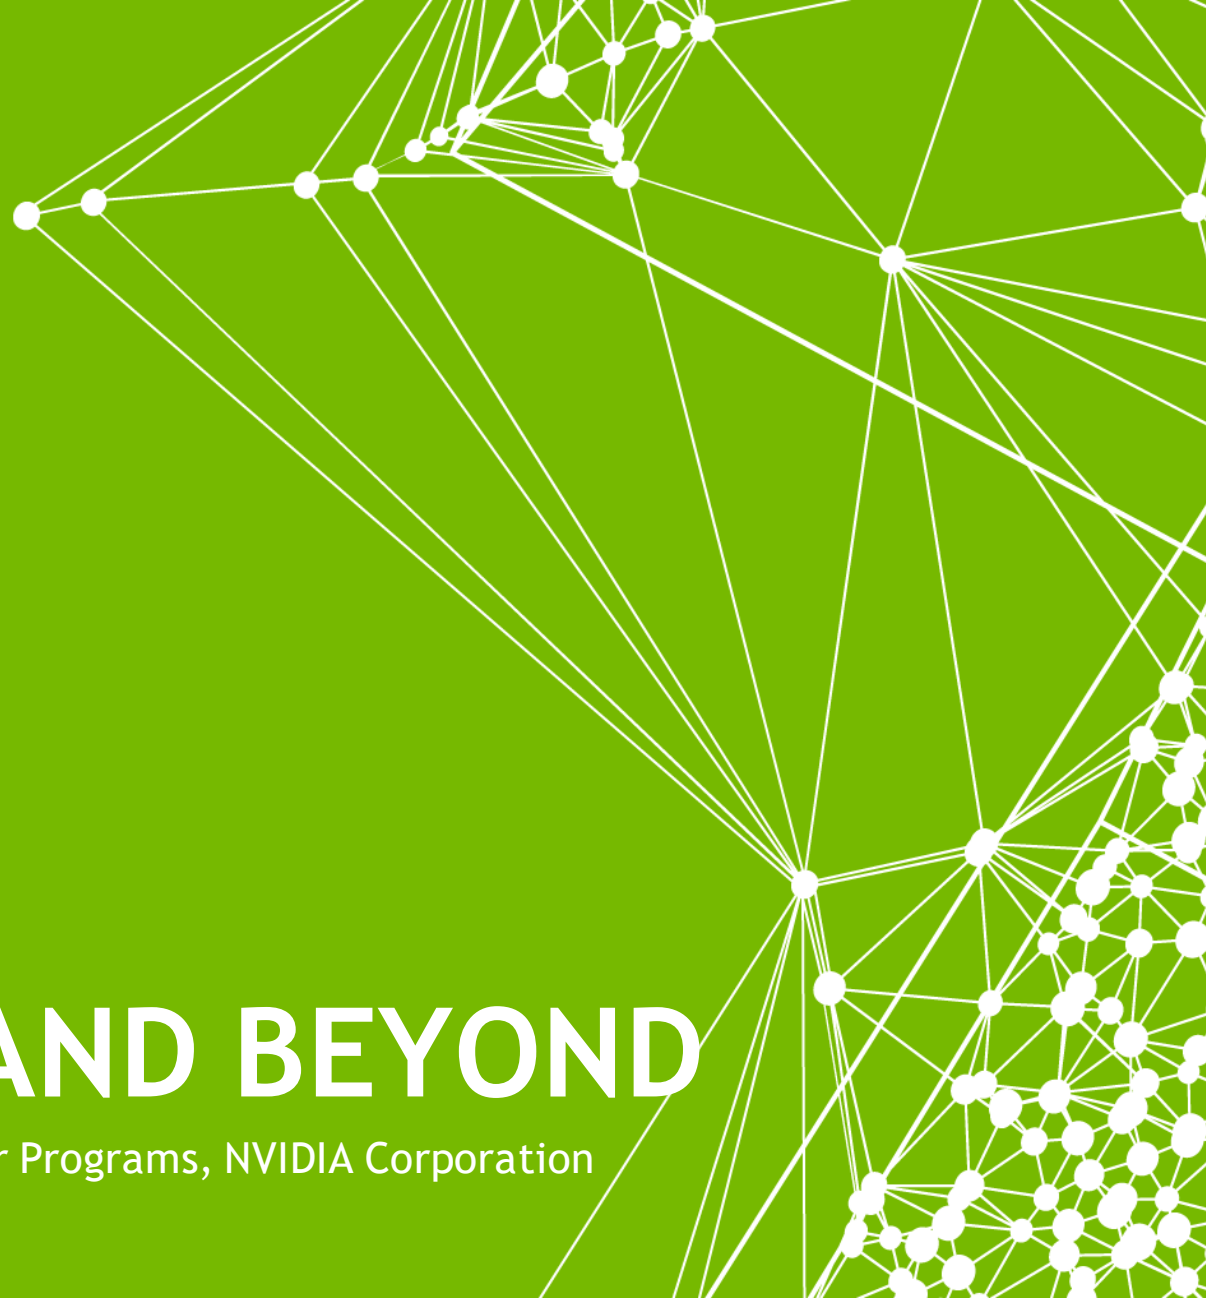

NVIDIA | DEEP LEARNING INSTITUTE

DEEP LEARNING AND BEYOND

Will Ramey, Sr. Director, Global Head of Developer Programs, NVIDIA Corporation

ACCELERATED DATA SCIENCE

DATA ANALYTICS

Extracting insights from big data

MACHINE LEARNING

Learning from examples in the data

DEEP LEARNING

Automating feature engineering

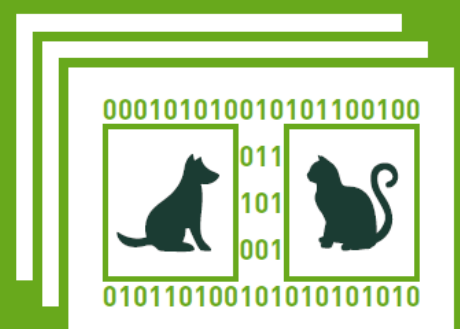

GPU-ACCELERATED DATA SCIENCE

Use Cases in Every Industry

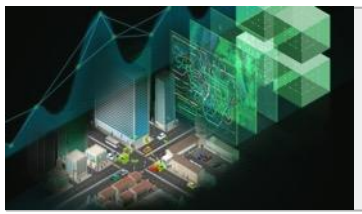

CONSUMER INTERNET

Ad Personalization
Click Through Rate Optimization
Churn Reduction

FINANCIAL SERVICES

Claim Fraud
Customer Service Chatbots/Routing
Risk Evaluation

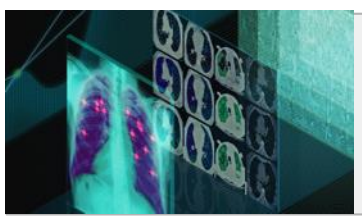

HEALTHCARE

Improve Clinical Care
Drive Operational Efficiency
Speed Up Drug Discovery

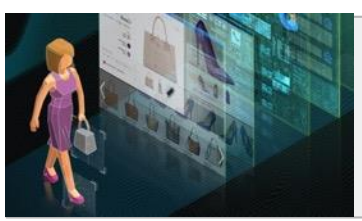

RETAIL

Supply Chain & Inventory Management
Price Management / Markdown Optimization
Promotion Prioritization And Ad Targeting

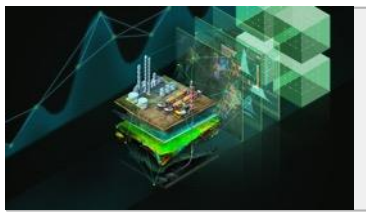

OIL & GAS

Sensor Data Tag Mapping
Anomaly Detection
Robust Fault Prediction

MANUFACTURING

Remaining Useful Life Estimation
Failure Prediction
Demand Forecasting

TELECOM

Detect Network/Security Anomalies
Forecasting Network Performance
Network Resource Optimization (SON)

AUTOMOTIVE

Personalization & Intelligent Customer Interactions
Connected Vehicle Predictive Maintenance
Forecasting, Demand, & Capacity Planning

BEYOND DEEP LEARNING

Opportunities to Accelerate Data Science

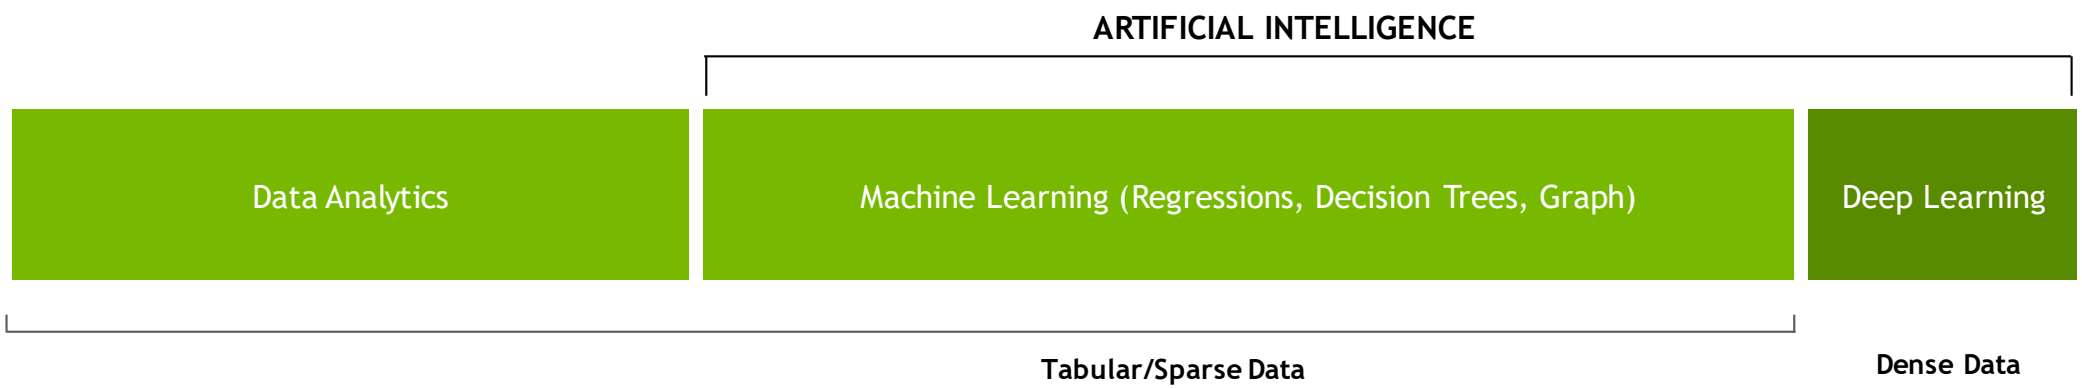

2.2 exabytes (2.2B GB) of data created daily – McKinsey

$260B annual revenue by 2022 for big data and business analytics – IDC

THE EXPANDING UNIVERSE OF MODERN AI

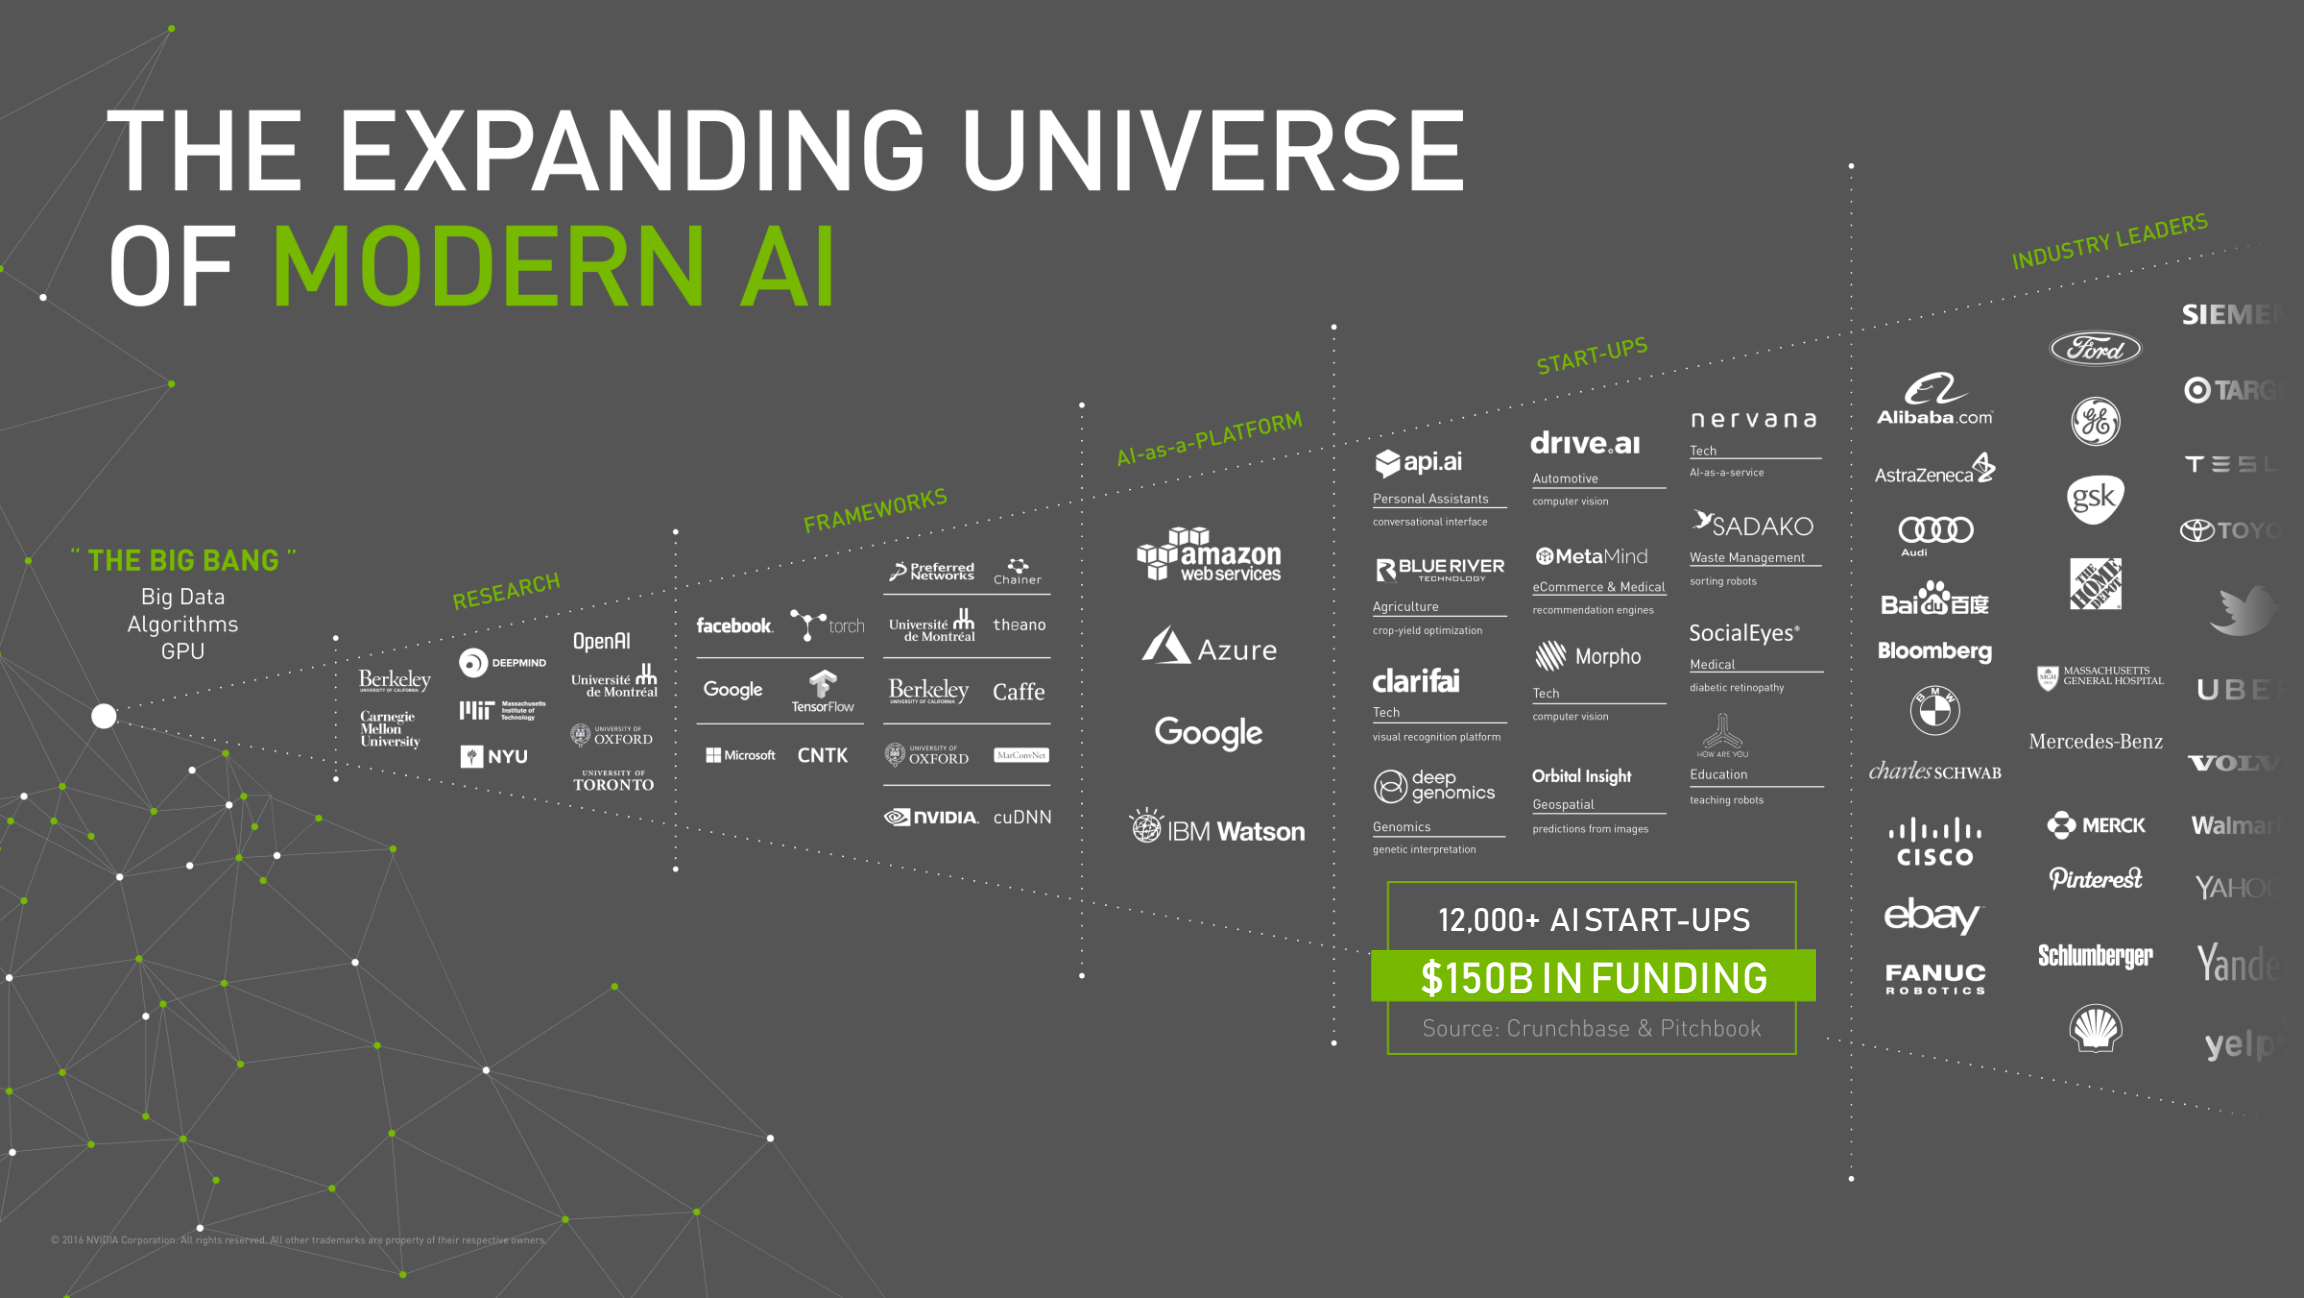

WHAT PROBLEM ARE YOU SOLVING?

Defining the AI/DL task

INPUTS	BUSINESS QUESTIONS	AI / DL TASK	EXAMPLE OUTPUTS		
			HEALTHCARE	RETAIL	FINANCE
Text Data Images Video Audio	Is “it” **present** or not?	Detection	Cancer Detection	Targeted Ads	Cybersecurity
	What **type** of thing is “it”?	Classification	Image Classification	Basket Analysis	Credit Scoring
	To what **extent** is “it” present?	Segmentation	Tumor Size / Shape Analysis	Build 360° Customer View	Credit Risk Analysis
	What is the likely **outcome**?	Prediction	Survivability Prediction	Sentiment & Behavior Recognition	Fraud Detection
	What will likely satisfy the objective?	Recommendations	Therapy Recommendation	Recommendation Engine	Algorithmic Trading

CAMBRIAN EXPLOSION

Convolutional Networks

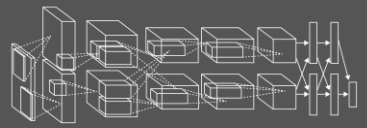


Encoder/Decoder


ReLu


BatchNorm


Concat


Dropout


Pooling

Recurrent Networks

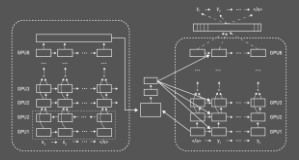


LSTM


GRU


Beam Search


WaveNet


CTC


Attention

Generative Adversarial Networks

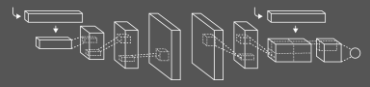


3D-GAN


MedGAN


Conditional GAN

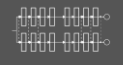
Coupled GAN

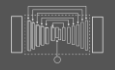
Speech Enhancement GAN

Reinforcement Learning

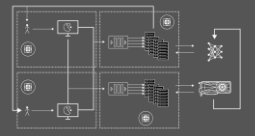


DQN


Simulation

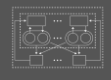
DDPG

New Species

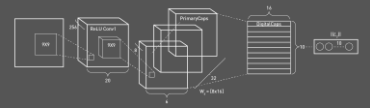

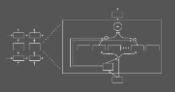
Mixture of Experts


Neural Collaborative Filtering


Block Sparse LSTM

DEEP LEARNING APPLICATION DEVELOPMENT

DEEP LEARNING APPLICATION DEVELOPMENT

Untrained
Neural Network
Model

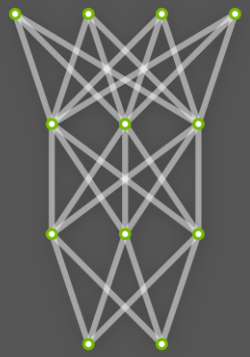

TRAINING
DATASET

DEEP LEARNING APPLICATION DEVELOPMENT

TRAINING

Learning a new capability
from existing data

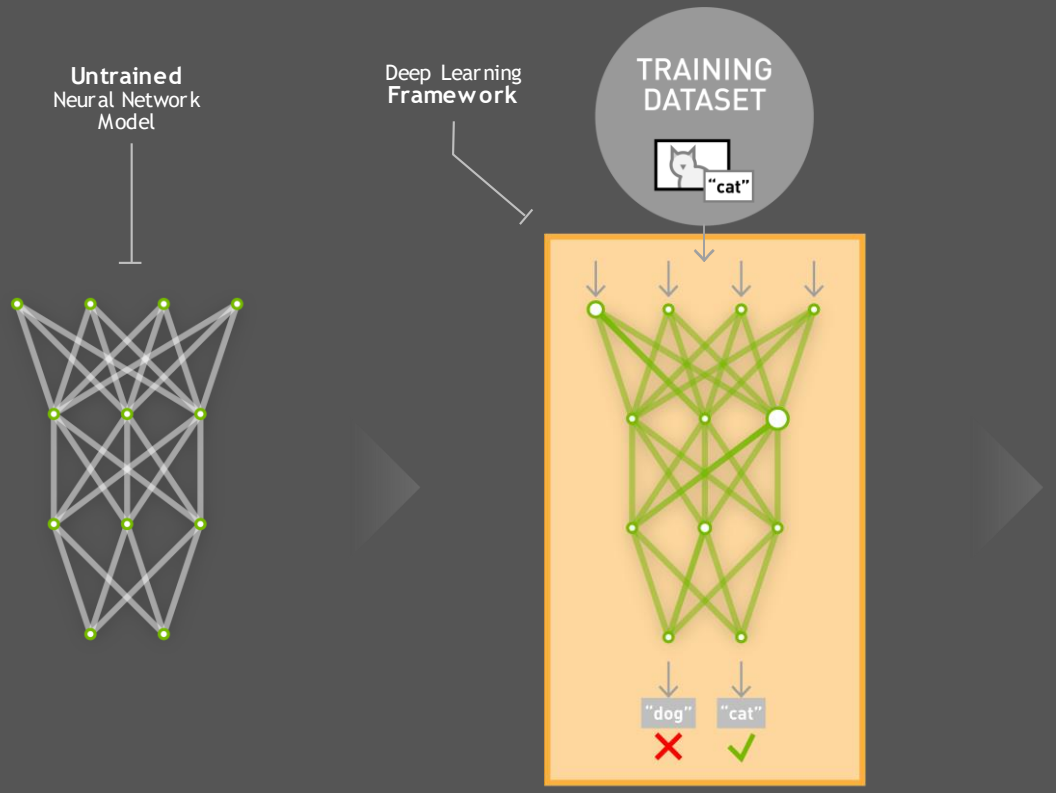

DEEP LEARNING APPLICATION DEVELOPMENT

TRAINING

Learning a new capability from existing data

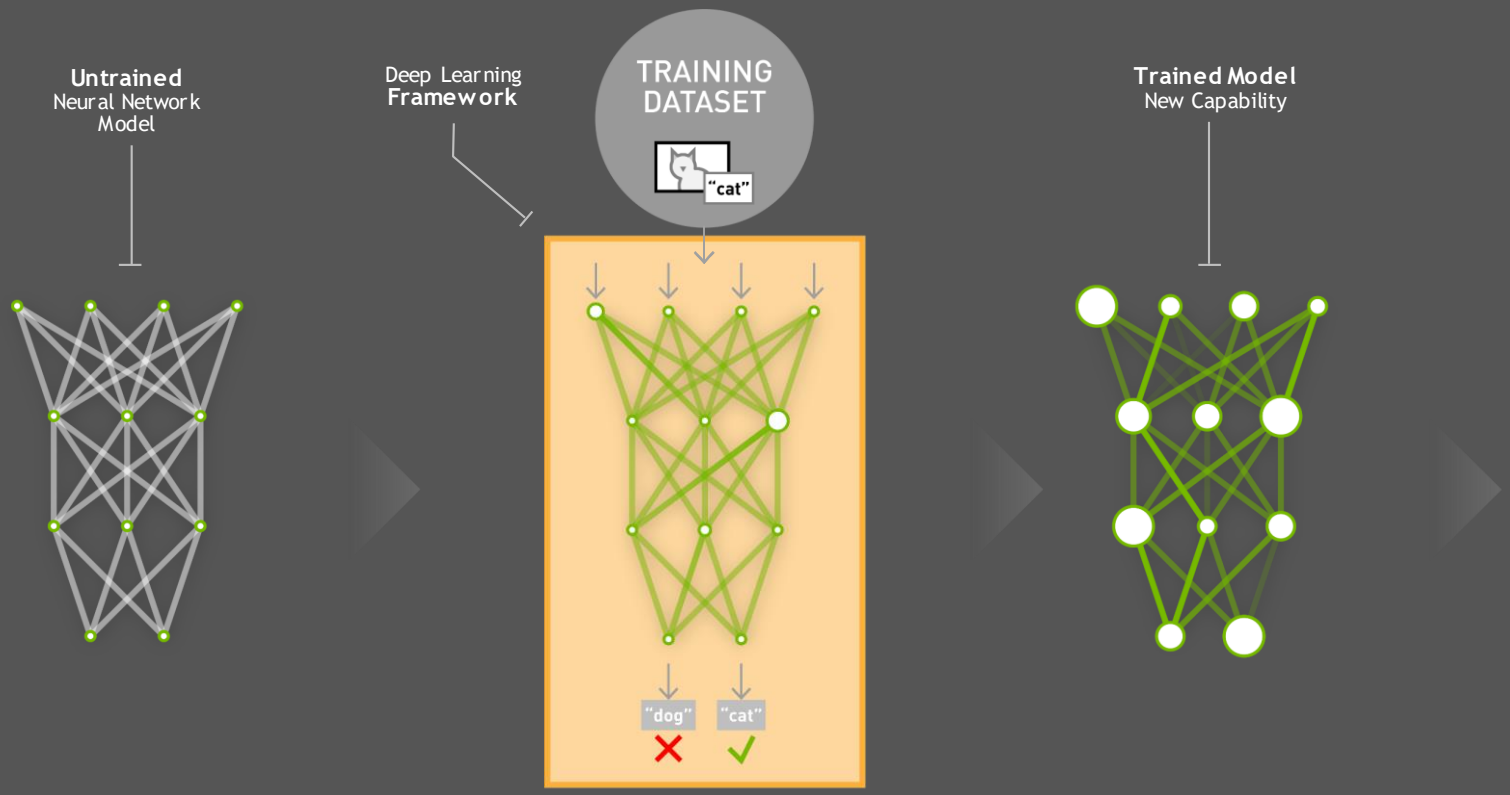

DEEP LEARNING APPLICATION DEVELOPMENT

TRAINING
Learning a new capability
from existing data

Untrained
Neural Network
Model

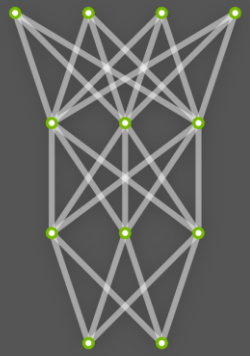

Deep Learning
Framework

TRAINING
DATASET

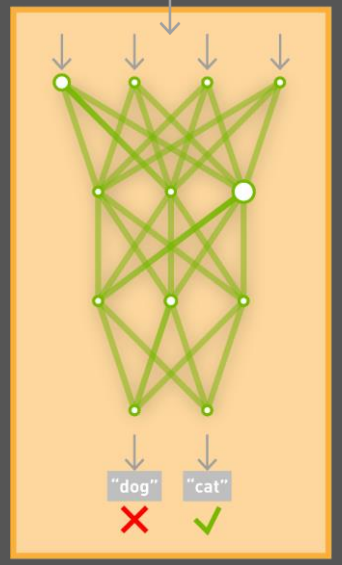

Trained Model
New Capability

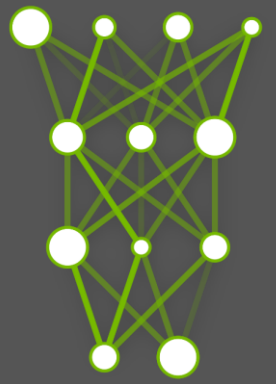

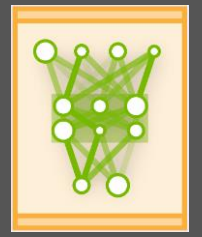

Trained Model
Optimized for
Performance

DEEP LEARNING APPLICATION DEVELOPMENT

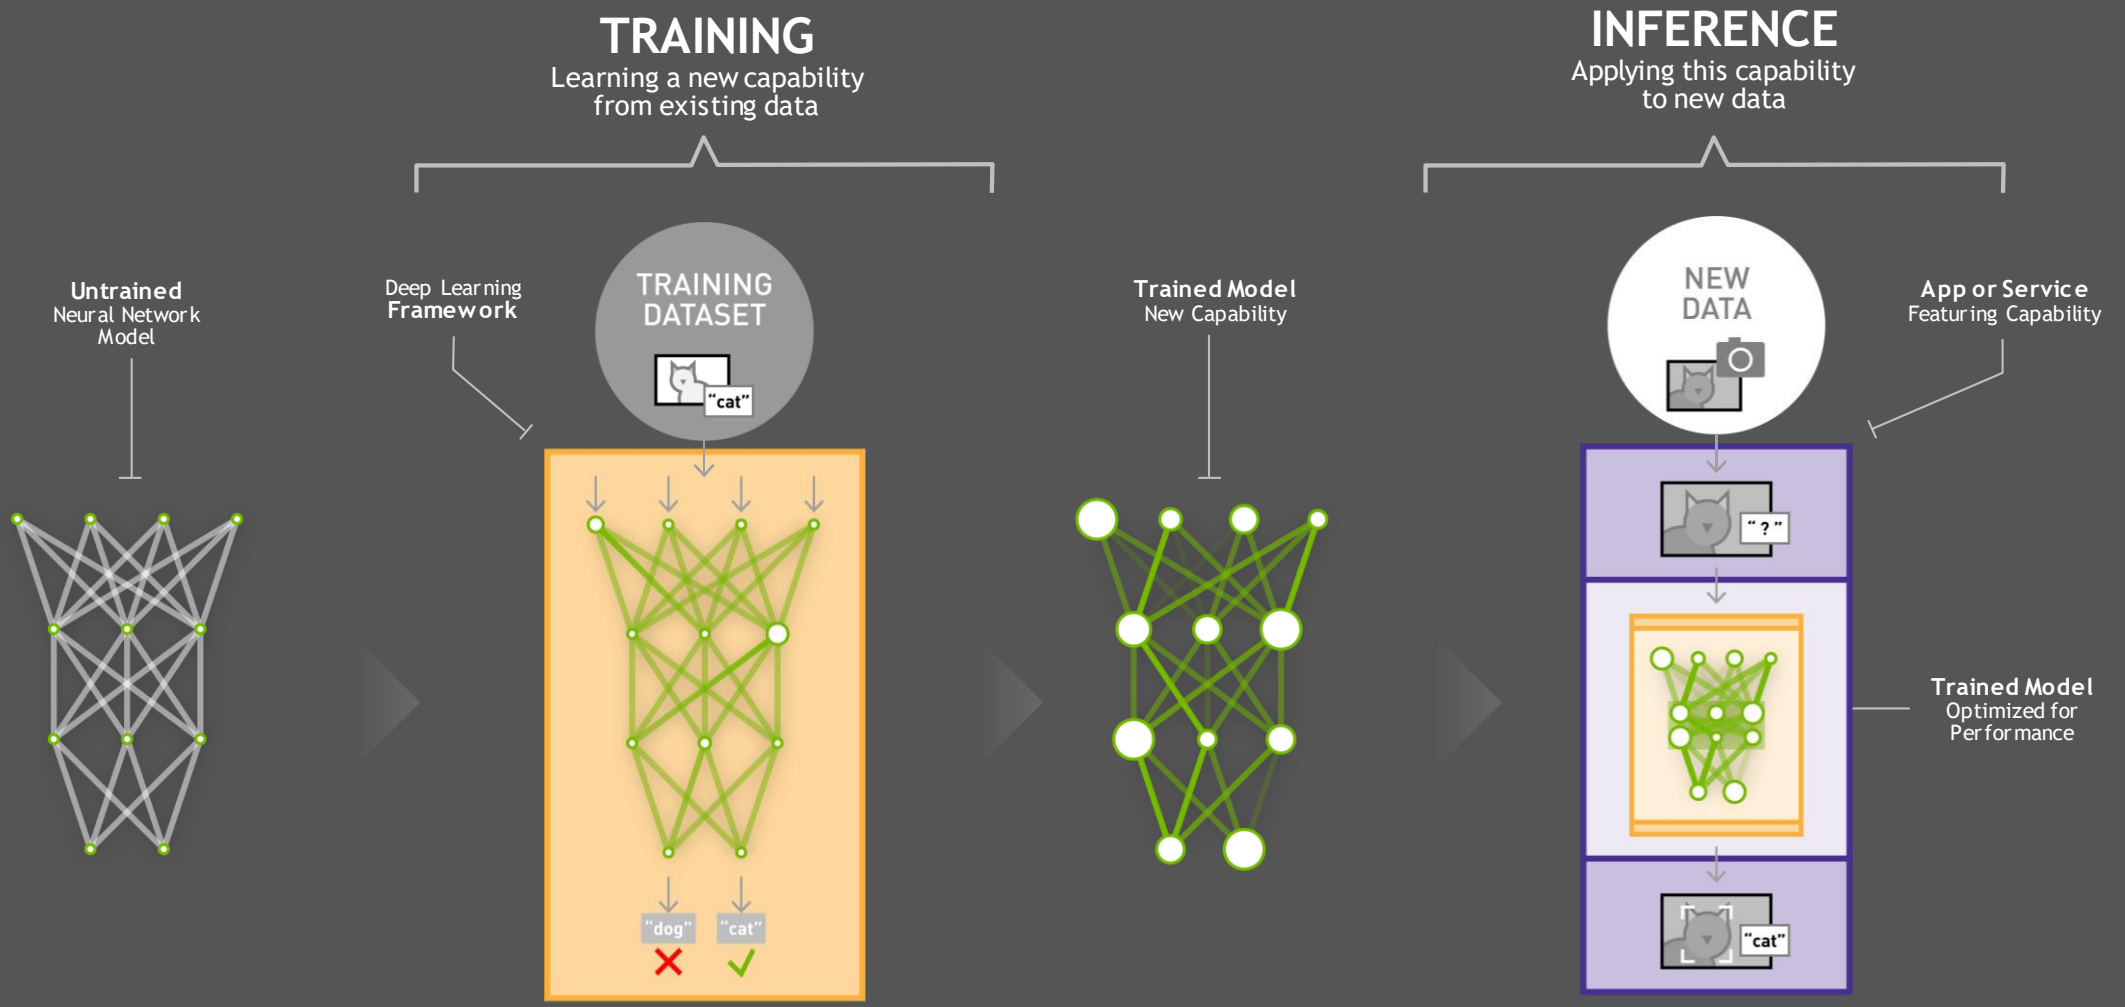

CHALLENGES

DEEP LEARNING NEEDS	WHY
Data Scientists	New computing model
Latest Algorithms	Rapid evolving
Fast Training	Impossible -> Practical
Deployment Platforms	Must be available everywhere

NVIDIA DEEP LEARNING INSTITUTE

Hands-on Training for Data Scientists and Software Engineers

Helping the world to solve challenging problems using AI and deep learning

On-site workshops and online courses presented by certified instructors

Covering complete workflows for proven application use cases

Fundamentals, Autonomous Vehicles, Finance, Healthcare, Video Analytics, IoT/Robotics, …

www.nvidia.com/dli

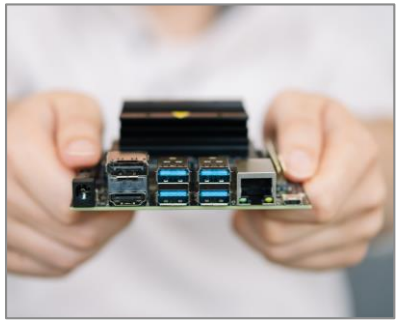

Getting Started with AI on Jetson Nano

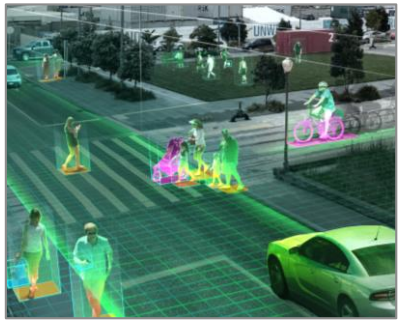

Deep Learning for Intelligent Video Analytics

Deep Learning for Robotics

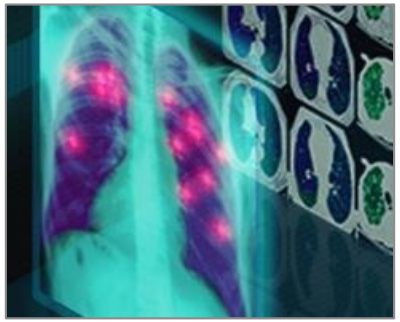

Deep Learning for Healthcare Image Analysis

Fundamentals of Accelerated Data Science with RAPIDS

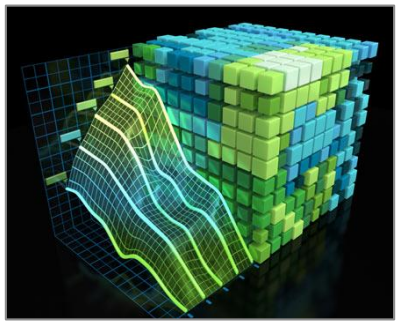

Fundamentals of Accelerated Computing with CUDA Python

NVIDIA GTC

THE PREMIER AI CONFERENCE

MARCH 22-26, 2020 | Silicon Valley | #GTC20

ADVANCE YOUR DEEP LEARNING KNOWLEDGE AT GTC

Don't miss the world's most important event for GPU developers

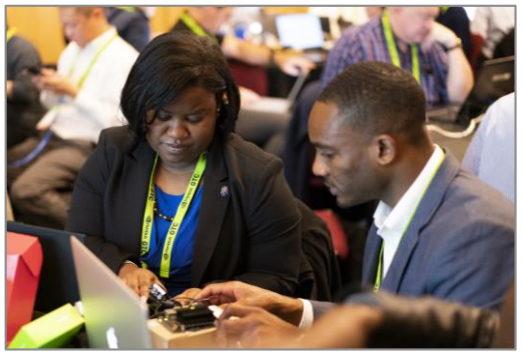

DEEP LEARNING INSTITUTE

KEYNOTES

SHOW FLOOR

PANELS

JOIN US AT GTC 2020 | USE VIP CODE NVWRAMEY FOR 25% OFF

Silicon Valley | WORKSHOPS MARCH 22nd | CONFERENCE MARCH 23rd -26th

NVIDIA INCEPTION: ACCELERATING AI STARTUPS

BENEFITS

AI Expertise

NVIDIA DEEP LEARNING INSTITUTE

Go-To-Market Support

Technology Access

aws

MEMBERSHIP

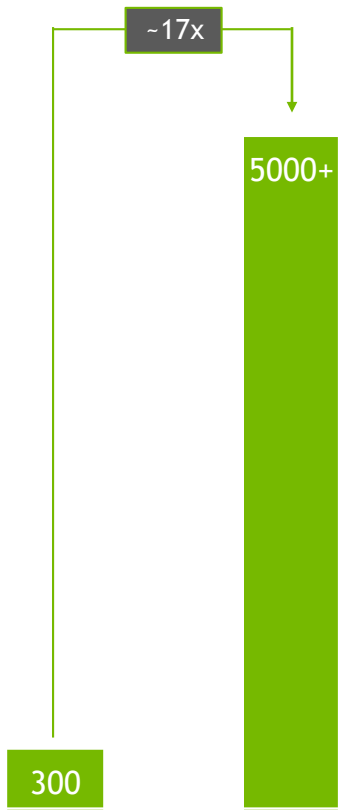

IMPACT

Showcasing Innovation

AIFI · athelas · Caption Health · deepinstinct · focal systems · Genetesis · KINEMA SYSTEMS · SMARTVID.IO · SUBTLE MEDICAL · meero · omni·sci · Skydio

Creating a Global Community

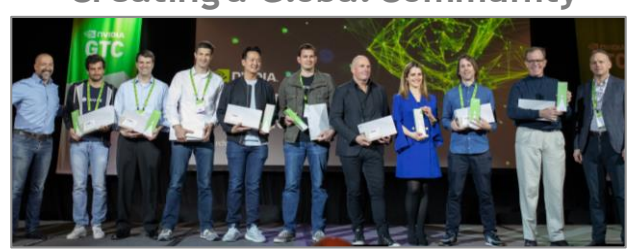

www.nvidia.com/inception

NVIDIA DEEP LEARNING SOFTWARE PLATFORM

TRAINING

Training Data

Data Management

Training

Model Assessment

Trained Neural Network

Chainer | Microsoft Cognitive Toolkit | mxnet | PyTorch | TensorFlow

INFERENCE

Data Center — GRE + TensorRT

Embedded — JETPACK SDK

Automotive — DriveWorks SDK

NVIDIA DEEP LEARNING SDK and CUDA

cuDNN | NCCL (GPU0, GPU1, GPU2, GPU3) | cuBLAS | cuSPARSE | TensorRT | DeepStream SDK (NVDEC, Inference, NVENC)

developer.nvidia.com/deep-learning-software

NGC

Simple Access to a Comprehensive Catalog of GPU-accelerated Software

Discover 30 GPU-Accelerated Containers
Deep learning, third-party managed HPC applications,
NVIDIA HPC visualization tools, and partner applications

Innovate in Minutes, Not Weeks
Get up and running quickly and reduce complexity

Access from Anywhere
Use on PCs with NVIDIA Volta or Pascal™ architecture GPUs,
NVIDIA DGX Systems, and supported cloud providers

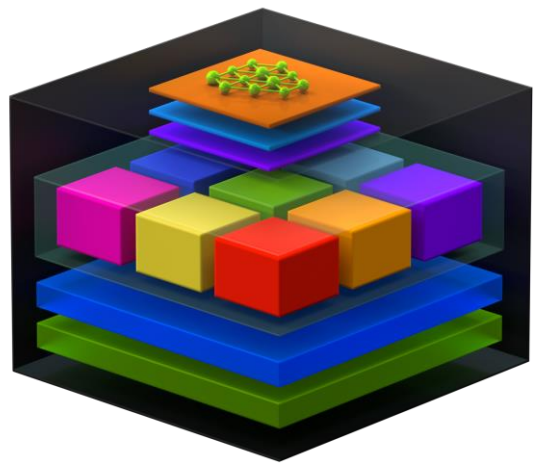

ngc.nvidia.com

NVIDIA EGX PLATFORM

AI Scalability from Edge to Data Center

Securely deploy and manage containerized AI frameworks and applications, including NVIDIA TensorRT™, TensorRT Inference Server, and DeepStream

Includes Kubernetes plug-in, container runtime, NVIDIA drivers, and GPU monitoring

Optimized for NGC-Ready Edge Systems, from Jetson Nano to GPU servers in the cloud or datacenter

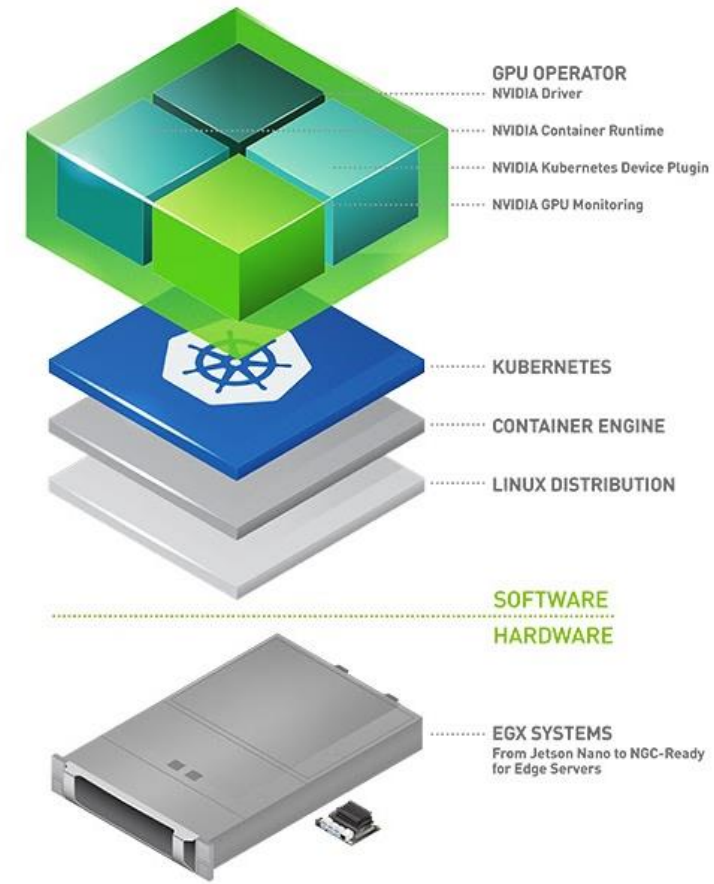

www.nvidia.com/egx

END-TO-END PRODUCT FAMILY

HPC / TRAINING

NVIDIA GPU CLOUD

INFERENCE

FULLY INTEGRATED AI SYSTEMS

DGX-1

DGX-2

DESKTOP

TITAN

WORKSTATION

DGX Station

VIRTUAL WS

Virtual GPU

DATA CENTER

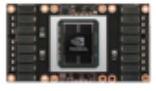

Tesla V100

SERVER PLATFORM

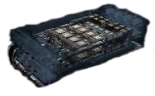

HGX1/HGX2

DATA CENTER

Tesla T4

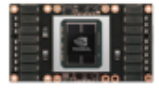

Tesla V100

AUTOMOTIVE

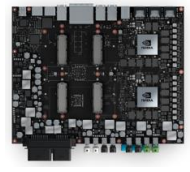

Drive AGX Pegasus

EMBEDDED

Jetson AGX Xavier

SOLUTIONS

DEEP LEARNING NEEDS	SOLUTIONS
Data Scientists	Deep Learning Institute, GTC
Latest Algorithms	GPU Accelerated Frameworks, DL SDK
Fast Training	DGX, V100, TITAN V
Deployment Platforms	TensorRT, NGC, EGX, V100/T4, Drive AGX, Jetson AGX

READY TO GET STARTED?

Project Checklist

What problem are you solving,
what are the AI/DL tasks?

What data do you have/need,
how is it labeled?

Which tools & environment will you use?

On what platform(s) will you train and deploy?

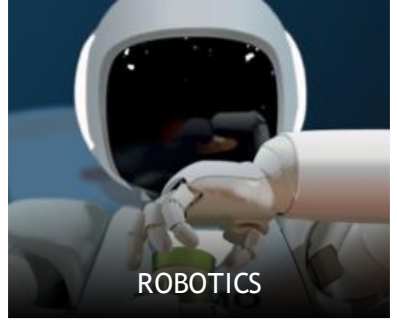

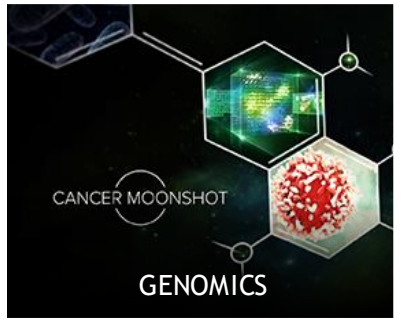

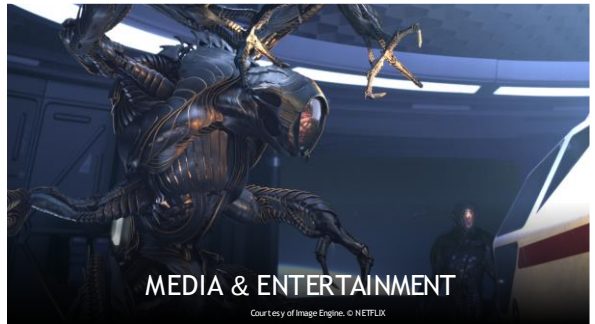

WHAT'S NEXT?

Review examples of AI in action	news.developer.nvidia.com
Listen to the NVIDIA AI Podcast	blogs.nvidia.com/ai-podcast
Experience DLI online courses	www.nvidia.com/dli
Register for GTC near you	www.nvidia.com/gtc

THANK YOU

NVIDIA | DEEP LEARNING INSTITUTE

www.nvidia.com/dli

NVIDIA

GTC GPU TECHNOLOGY CONFERENCE

THE LATEST DEEP LEARNING DEVELOPER TOOLS

March 22 | Full-Day Workshops
March 23 - 26 | Conference & Training

Get the hands-on experience you need to transform the future of AI, high-performance computing and more with NVIDIA's Deep Learning Institute (DLI).

Register for GTC 2020 to earn certification in full-day workshops, join instructor-led sessions, and start self-paced training.

www.nvidia.com/en-us/gtc/sessions/training/